

Deep neural network

Kyung-Ah Sohn




Ajou University

Contents


- Perceptron
- Learning XOR
- NN architecture design and computation
- Training NN: Back-propagation
- Regularization

History of neural networks

- **First generation (1958~): perceptron** (F. Rosenblatt, 1958)
 - Criticized by Marvin Minsky about XOR problem
- **Second generation (1986~) : multilayer perceptrons**
 - Trained by back-propagating error signal (1986)
 - Mostly shallow network with 1 hidden layer
- **Third generation (2006~): deep learning**
 - Deep belief nets (Hinton, 2006)
 - Deep neural network (DNN), convolutional neural network (CNN), ...



Recall: Logistic regression



Model

linear decision boundary

$$\log \frac{p(y = 1|\mathbf{x})}{1 - p(y = 1|\mathbf{x})} = \mathbf{w}^T \mathbf{x} + b$$

$$\rightarrow p(y = 1|\mathbf{x}) = \frac{1}{1 + e^{-(\mathbf{w}^T \mathbf{x} + b)}}$$

Input
features

$$(x_1, x_2)$$

$$\mathbf{w} = (w_1, w_2), b$$


$$t = w_1 x_1 + w_2 x_2 + b$$

Output

$$g(t) = \frac{1}{1 + e^{-t}}$$

$g(\cdot)$


Input
features



Output

Perceptron

- Perceptron with 2 input features




$$y = \begin{cases} 0 & (w_1x_1 + w_2x_2 \leq \theta) \\ 1 & (w_1x_1 + w_2x_2 > \theta) \end{cases}$$

threshold


Perceptron

Representation with bias




$$y = h(b + w_1x_1 + w_2x_2)$$

$$h(x) = \begin{cases} 0 & (x \leq 0) \\ 1 & (x > 0) \end{cases}$$





Perceptron





Perceptron: Boolean operation

x_1	x_2	y
0	0	0
0	1	0
1	0	0
1	1	1




x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	1




XOR is not linearly separable

Cannot be solved by a simple perceptron
But can be separated using AND and OR


x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	0




Building XOR from AND and OR




Artificial neuron



Multiple neurons form a **layer**





Multiple layers form a **network**



Artificial neural network

General Structure




We can form **non-linear functions** by composing stages of processing

Training ANN means learning the weights of the neurons



Network topology

- The number of layers
- The number of nodes within each layer
- Whether information in the network is allowed to travel backward



Direction of information travel

- Feed-forward network
 - Signal is fed in one direction
- Recurrent network (or feedback network)
 - Allows signal to travel in both directions using loops
 - Addition of a short term memory




Neural network mathematics

Neural network: input / output transformation

$$y_{out} = F(x, W)$$

W is the matrix of all weight vectors.




$$s = w_0 + w_1x_1 + w_2x_2 + \cdots + w_Mx_M = \mathbf{w}^T \mathbf{x}$$

Diagram illustrating the linear combination of input features and weights to produce the sum s . The input features are x_1, x_2, \dots, x_M , and the weights are w_1, w_2, \dots, w_M . The bias term w_0 is also included. The sum s is then passed through a non-linearity function $h = \sigma(s)$ to produce the output feature h .

$$\text{output feature} = \sigma(\text{sum})$$

Diagram illustrating the vector representation of the neuron model. The input features are represented as a vector \mathbf{x} , and the weights are represented as a vector \mathbf{w} . The sum s is calculated as the dot product $\mathbf{w}^T \mathbf{x}$. The output feature h is then calculated as $h = \sigma(s)$, where σ is the non-linearity function.

Layer: parallelized weighted sum and non-linearity



one sum per weight vector $s_j = \mathbf{w}_j^T \mathbf{x} \longrightarrow \mathbf{s} = \mathbf{W}^T \mathbf{x}$ vector of sums from weight matrix

$$\mathbf{h} = \sigma(\mathbf{s})$$




A diagram showing the matrix representation of the layer operation. It consists of two equations. The top equation is:

$$\text{sum} \begin{bmatrix} \vdots \end{bmatrix} = \begin{bmatrix} \text{weights} & \text{input features} \end{bmatrix}$$
 where the 'sum' is a column vector, the 'weights' is a matrix, and the 'input features' is a column vector. The bottom equation is:

$$\text{output feature} \begin{bmatrix} \vdots \end{bmatrix} = \sigma \left(\begin{bmatrix} \vdots \end{bmatrix} \right)$$
 where the 'output feature' is a column vector, σ is the non-linearity function, and the argument is a column vector labeled 'sum'.

Network: sequence of parallelized weighted sums and non-linearities



$$\begin{array}{ll} \text{1st layer} & \text{2nd layer} \\ \mathbf{s}^{(1)} = \mathbf{W}^{(1)} \mathbf{x}^{(0)} & \mathbf{s}^{(2)} = \mathbf{W}^{(2)} \mathbf{x}^{(1)} \\ \mathbf{x}^{(1)} = \sigma(\mathbf{s}^{(1)}) & \mathbf{x}^{(2)} = \sigma(\mathbf{s}^{(2)}) \end{array} \quad \dots$$


...

↓

$$\text{output} = \sigma(\dots \sigma(\text{2nd weights} \sigma(\text{1st weights} \text{input})) \dots)$$


Activation functions

- The mechanism by which the artificial neuron processes information and passes it throughout the network
- Create non-linearity
- Unit step function (threshold activation function)
 - Not continuous
 - Returns either 0 or 1 (discrete)
 - Non-linear function



Activation functions

- Sigmoid function
 - The values of logistic function range from 0 to 1 (continuous)
 - Differentiable (mathematically)
 - Non-linear
 - The most commonly used in traditional ANN




Activation functions


- **tanh**: takes a real-valued input and squashes it to the range $[-1, 1]$
- **ReLU**: ReLU stands for Rectified Linear Unit. It takes a real-valued input and thresholds it at zero (replace negative values with zero)

$$f(x) = \max(0, x)$$


$$h(x) = \begin{cases} x & (x > 0) \\ 0 & (x \leq 0) \end{cases}$$



Sigmoid



tanh




ReLU


Activation functions

- Traditionally

logistic sigmoid



hyperbolic tangent (tanh)




***sat*urating**


*derivative goes to
zero at $+\infty$ and $-\infty$*

- More recently

rectified linear unit (ReLU)




leaky ReLU



softplus



exponential linear unit (ELU)




***non-sat*urating**

*non-zero derivative
at $+\infty$ and/or $-\infty$*

Activation at output layer


For regression

- Identity function



For classification

- Softmax function



$$y_k = \frac{\exp(a_k)}{\sum_{i=1}^n \exp(a_i)}$$

Loss function

ex: MNIST classification

- Measure to calculate how 'bad' the NN is
- e.g. MNIST classification

```
>>> y = [0.1, 0.05, 0.6, 0.0, 0.05, 0.1, 0.0, 0.1, 0.0, 0.0] : Output from NN  
>>> t = [0, 0, 1, 0, 0, 0, 0, 0, 0, 0] : True label
```

– Mean squared error (MSE)

$$E = \frac{1}{2} \sum_k (y_k - t_k)^2$$

– Cross entropy

$$E = - \sum_k t_k \log y_k$$

Loss function

- Suppose the neural network's computed outputs and the target values are as follows

NN1

computed	targets	correct?
0.3 0.3 0.4	0 0 1 (democrat)	yes
0.3 0.4 0.3	0 1 0 (republican)	yes
0.1 0.2 0.7	1 0 0 (other)	no

NN2

computed	targets	correct?
0.1 0.2 0.7	0 0 1 (democrat)	yes
0.1 0.7 0.2	0 1 0 (republican)	yes
0.3 0.4 0.3	1 0 0 (other)	no


	NN1	NN2
Classification error	0.33	0.33
Avg. Cross-Entropy	1 st sample: $-(0 \cdot \log 0.3 + 0 \cdot \log 0.3 + 1 \cdot \log 0.4) = -\log 0.4$ $\rightarrow \text{ACE} = (-\log 0.4 - \log 0.4 - \log 0.1) / 3 = 1.38$	0.64
MSE	$(0.54 + 0.54 + 1.34) / 3 = 0.81$	0.34

Learning as optimization

- To learn the weights, we need the **derivative** of the loss w.r.t the weights

How should the weight be updated to decrease the loss?

$$w = w - \alpha \frac{\partial \mathcal{L}}{\partial w}$$



With multiple weights

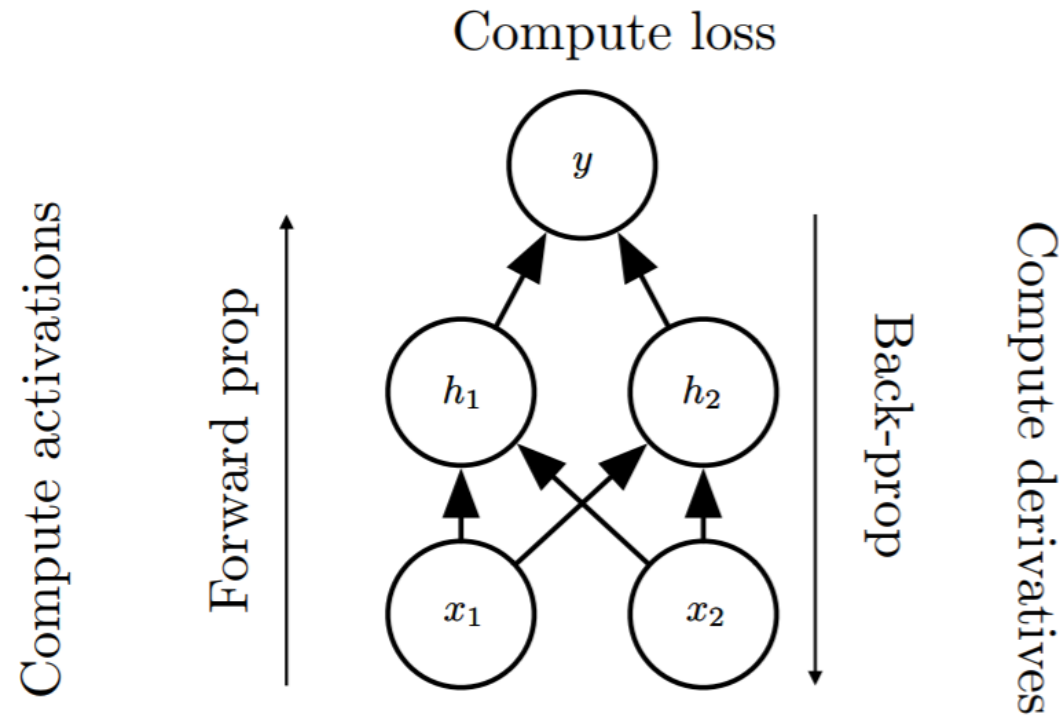
$$\mathbf{w} = \mathbf{w} - \alpha \nabla_{\mathbf{w}} \mathcal{L}$$

Learning rates (pointing to α)

gradients (pointing to $\nabla_{\mathbf{w}} \mathcal{L}$)

Stochastic Gradient Descent (SGD)

Simple Back-Propagation example



Backpropagation

A neural net defines a function of composed operations

$$f_L(\mathbf{w}_L, f_{L-1}(\mathbf{w}_{L-1}, \dots f_1(\mathbf{w}_1, \mathbf{x}) \dots))$$

and the loss L is a function of the network output

→ Use [chain rule](#) to calculate gradients

chain rule example

$$y = w_2 e^{w_1 x}$$

input x

output y

parameters w_1, w_2

evaluate parameter derivatives: $\frac{\partial y}{\partial w_1}, \frac{\partial y}{\partial w_2}$

define

$$v \equiv e^{w_1 x} \longrightarrow y = w_2 v$$

$$u \equiv w_1 x \longrightarrow v = e^u$$

then $\frac{\partial y}{\partial w_2} = v = e^{w_1 x}$

$\frac{\partial y}{\partial w_1} = \boxed{\frac{\partial y}{\partial v} \frac{\partial v}{\partial u} \frac{\partial u}{\partial w_1}} = w_2 \cdot e^{w_1 x} \cdot x$


chain rule

Backpropagation


- Recall

1st layer	2nd layer	...	Loss
$\mathbf{s}^{(1)} = \mathbf{W}^{(1)} \mathbf{x}^{(0)}$	$\mathbf{s}^{(2)} = \mathbf{W}^{(2)} \mathbf{x}^{(1)}$		\mathcal{L}
$\mathbf{x}^{(1)} = \sigma(\mathbf{s}^{(1)})$	$\mathbf{x}^{(2)} = \sigma(\mathbf{s}^{(2)})$		

- To determine the chain rule ordering, we draw the dependency graph



Backpropagation



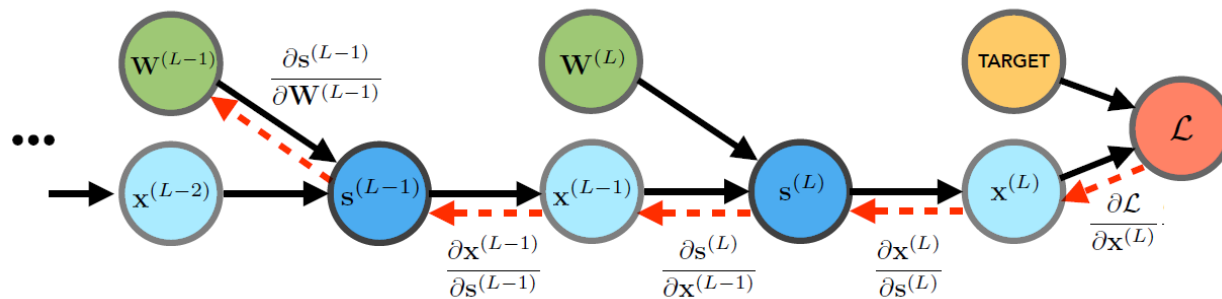
$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(L)}} = \frac{\partial \mathcal{L}}{\partial \mathbf{x}^{(L)}} \frac{\partial \mathbf{x}^{(L)}}{\partial \mathbf{s}^{(L)}} \frac{\partial \mathbf{s}^{(L)}}{\partial \mathbf{W}^{(L)}}$$

depends on the form of the loss derivative of the non-linearity $\frac{\partial}{\partial \mathbf{W}^{(L)}} (\mathbf{W}^{(L)} \mathbf{x}^{(L-1)}) = \mathbf{x}^{(L-1) \top}$

note $\nabla_{\mathbf{W}^{(L)}} \mathcal{L} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(L)}}$ is notational convention

Backpropagation


- Go back one more layer



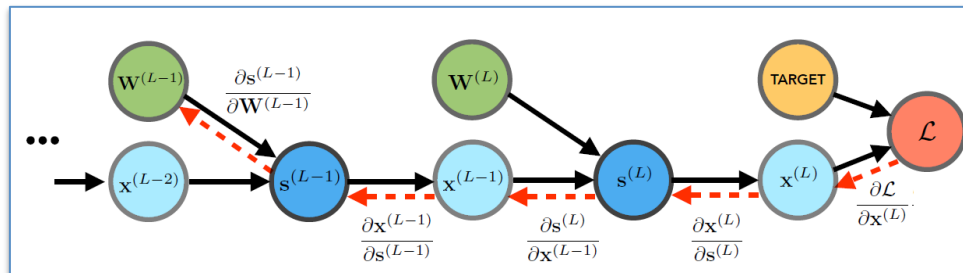
$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(L-1)}} = \frac{\partial \mathcal{L}}{\partial \mathbf{x}^{(L)}} \frac{\partial \mathbf{x}^{(L)}}{\partial \mathbf{s}^{(L)}} \frac{\partial \mathbf{s}^{(L)}}{\partial \mathbf{x}^{(L-1)}} \frac{\partial \mathbf{x}^{(L-1)}}{\partial \mathbf{s}^{(L-1)}} \frac{\partial \mathbf{s}^{(L-1)}}{\partial \mathbf{W}^{(L-1)}}$$

Backpropagation

- Some of the terms appear in both gradients
- e.g. we can reuse $\frac{\partial \mathcal{L}}{\partial \mathbf{s}^{(L)}}$




$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(L)}} = \frac{\partial \mathcal{L}}{\partial \mathbf{x}^{(L)}} \frac{\partial \mathbf{x}^{(L)}}{\partial \mathbf{s}^{(L)}} \frac{\partial \mathbf{s}^{(L)}}{\partial \mathbf{W}^{(L)}}$$



$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(L-1)}} = \frac{\partial \mathcal{L}}{\partial \mathbf{x}^{(L)}} \frac{\partial \mathbf{x}^{(L)}}{\partial \mathbf{s}^{(L)}} \frac{\partial \mathbf{s}^{(L)}}{\partial \mathbf{x}^{(L-1)}} \frac{\partial \mathbf{x}^{(L-1)}}{\partial \mathbf{s}^{(L-1)}} \frac{\partial \mathbf{s}^{(L-1)}}{\partial \mathbf{W}^{(L-1)}}$$


Backpropagation

- Update weights using gradients
- BP calculates the gradients via chain rule
- Gradient is propagated backward through the network




- Most deep learning software libraries automatically calculate gradients

Vanishing gradient



- In backprop, the product of many small terms goes to zero
- *Difficult to train very deep neural network with sigmoid*

ReLU



- In the positive region, ReLU does not saturate, preventing gradients from vanishing in deep networks
- In the negative region, ReLU saturates at zero, resulting in 'dead units', but in practice, this doesn't seem to be a problem
- Most commonly used in DNN

Training NN: summary



- 전제
 - 신경망의 가중치(weights)와 편향(bias)을 훈련데이터에 적응하도록 조정하는 과정을 '학습' 이라고 함
- 1단계- 미니배치
 - 훈련데이터 중 일부를 무작위로 가져옴. 이렇게 선별한 데이터를 미니배치라 하며, 그 미니배치의 손실(Loss) 값을 줄이는 것이 목표
- 2단계 – Gradient (기울기) 산출
 - 미니배치의 손실 함수 값을 줄이기 위해 각 가중치 매개변수(weight parameter)의 기울기를 구함. 기울기는 손실 함수의 값을 가장 작게 하는 방향을 제시함
- 3단계 – 매개변수 갱신
 - 가중치 매개변수를 기울기 방향으로 아주 조금 갱신함
- 4단계 – 반복
 - 1~3단계를 반복함.

$$E = \frac{1}{N} \sum_{i=1}^N (F(x_i; W) - y_i)^2$$

$$\Delta w_i^j = -c \cdot \frac{\partial E}{\partial w_i^j} (W)$$
$$w_i^{j, new} = w_i^j + \Delta w_i^j$$

Initialization

- Learning can be sensitive to weight initialization



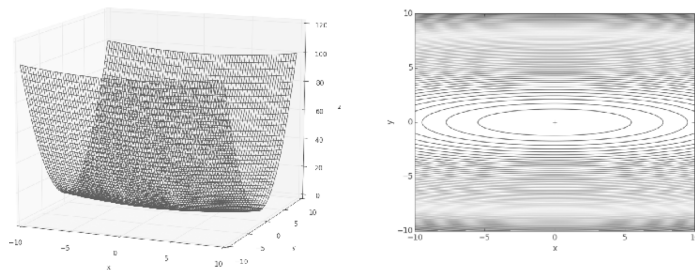
- Xavier* initialization (2010)
- He* initialization (2015)

Parameter update


- SGD (확률적 경사 하강법)
 - Simple, easy to implement
 - Inefficient sometimes

$$\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial L}{\partial \mathbf{W}}$$

$$f(x,y) = \frac{1}{20}x^2 + y^2$$



Optimization path




Parameter update

- Momentum (운동량)

$$\mathbf{W} \leftarrow \mathbf{W} + \mathbf{v}$$


$$\mathbf{v} \leftarrow \alpha \mathbf{v} - \eta \frac{\partial L}{\partial \mathbf{W}}$$

Optimization update path by momentum




Parameter update

AdaGrad



Adam



RMSprop, Adadelata, Adamax, ...

Overfitting

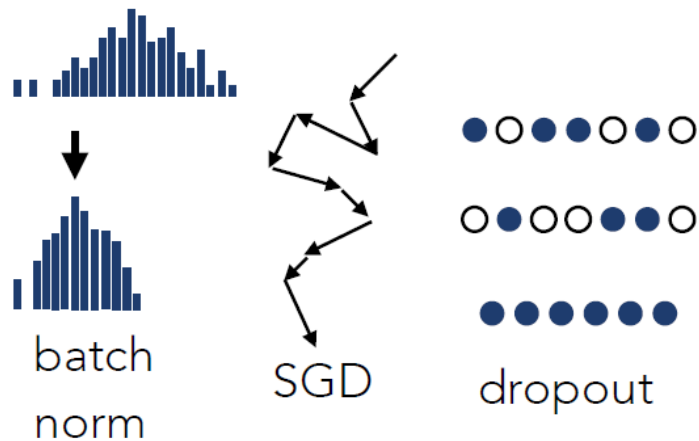
- Especially with
 - Complex models with many parameters
 - Small training data
- How to avoid
 - Weight decay (add L2 penalty on W to loss)
 - Dropout: randomly *drop*(remove) neurons during training

Regularization

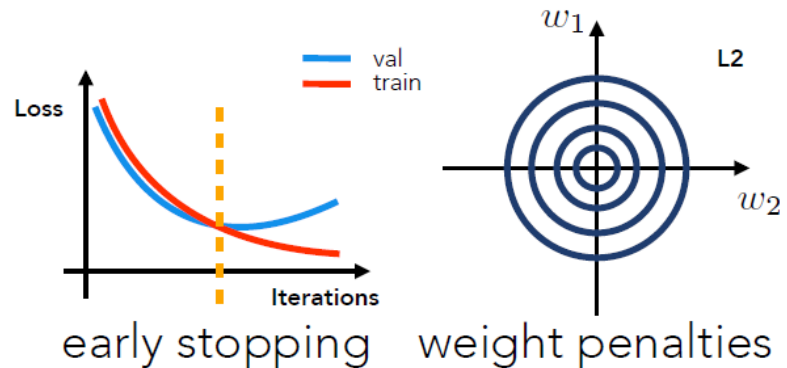
Regularization

- Regularization combats overfitting

stochasticity (uncertainty)




constraints



Batch normalization

- Splits the training dataset into small mini-**batches** that are used to calculate model error and update model coefficients.
- **Normalization** transform distribution into standard Normal

$$X_{\text{normal}} = \frac{X_{\text{original}} - \mu}{\sigma}$$




- **Batch norm**. normalizes each layer's activations according to the statistics of the batch
 - Results less sensitive to initialization

$$\mathbf{s}^{(\ell)} \leftarrow \gamma \frac{\mathbf{s}^{(\ell)} - \mu_{\mathcal{B}}}{\sigma_{\mathcal{B}}} + \beta$$

Hyperparameter tuning

- Hyperparameters
 - Number of nodes, Batch size
 - Learning rate (initial, decay schedule), regularization strength (L2 penalty, dropout)
 - Tuned over validation set




Hyperparameter tuning

- Search on log scale

```
learning_rate = 10 ** uniform(-6, 1)
```

- Random search vs. grid search



Grid

Random

- From coarse to fine ranges

Learning automatically from data



... End-to-end learning